

Measurement analysis of cacheability across multiple methods of grey-scaling

Niels de Waal (1698041), Jasper Smienk(1700502)

April 15, 2018

Contents

1	Target	3
2	Hypothesis	3
3	Method	3
4	Results	5
4.1	Raw data	5
4.2	Visualized	6
5	Processing	7
6	Conclusion	8
7	Evaluation	8
7.1	What went well	8
7.2	What went wrong	8
7.3	What could be done differently	8
8	Additional information	9
8.1	LSCPU	9

1 Target

This report will focus on the cacheability of 3 different methods of grey-scaling. These will be compared to each other and the default method for grey-scaling.

These results can help decide on which algorithm to use in a facial recognition project. More and more small systems are getting an on-cpu cache. This cache should be very useful in situations where a lot of data needs to be read and processed by the cpu. Which is the case for facial recognition, where every frame needs to be read from memory and processed by the cpu. In this report, there will be an investigation using multiple methods of grey-scaling and seeing how well these methods behave with respect to the cache.

2 Hypothesis

We can't say anything about how they will perform against the default implementation, as we don't know how it works.

We suspect there will be very little difference between the different methods when they will be examined on their own. It is unlikely that the results will change with or without the facial recognition enabled. We do expect to see a lot of data going to the L2 cache. This because the cpu pre-fetcher should see the sequential reads from the pixel buffer and try to cache data ahead of these reads.

3 Method

The binary will run one of the grey-scaling methods 10000 times, and we will run this binary 10 times, after which we will aggregate the results. We will use the *perf* tool for collecting the data. *Perf* is a linux tool often used for measuring kernel performance. In this test we will look at both L1-dcache and L2 cache. Seeing how a single pixel is quite small in terms of the size of the data (3 * 8 bits). It could be beneficial to know where most of the data ends up.

The test will be run twice for each method, once with the facial recognition, and once without.

We will make sure the tests are run on the same system and keep an eye out on the temperature to make sure it does not thermal-throttle.

Specifications of the used laptop:

CPU	Intel Core i5 6600K [1]
RAM	16GB DDR4 2400MHz
OS	Linux 4.15.15-1-ARCH
L1 DCache	32K
L2 Cache	256K

Used compiler optimisations:

- -std=c++14
- -g3
- -ggdb
- -g

- -O1

Used compiler warnings:

- -Wall
- -Wextra
- -pedantic-errors
- -Wfatal-errors
- -Wcast-align
- -Wmissing-declarations
- -Wredundant-decls
- -Wuninitialized
- -Wno-unused-parameter
- -Wno-missing-field-initializers

Perf command used:

```
$ perf stat -r 10 -d -e  
    l2_rqsts.demand_data_rd_hit,l2_rqsts.demand_data_rd_miss ./vissen
```

4 Results

4.1 Raw data

With facial recognition:

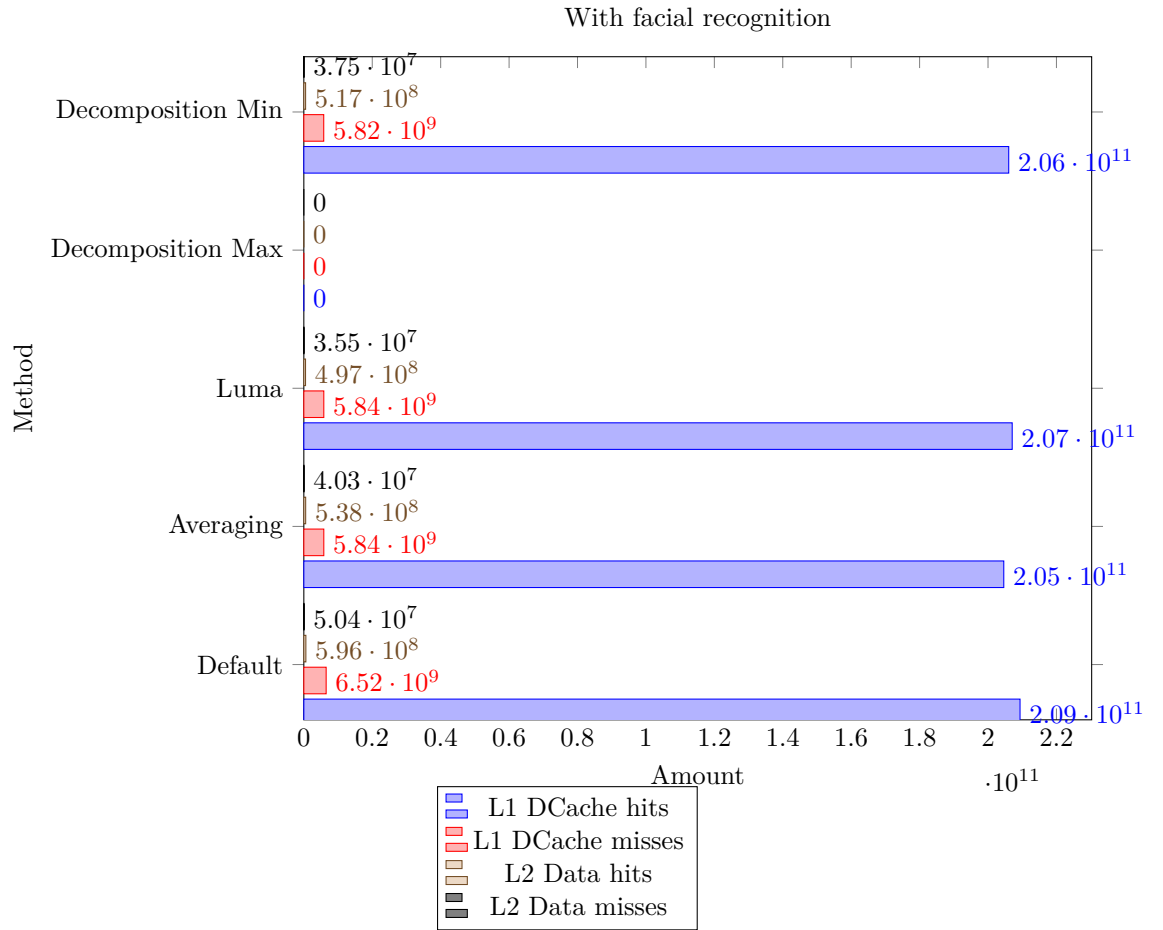
Method	L1 DCache hits	L1 DCache misses	L2 Data hits	L2 Data misses
Default	209359655680	6524946521	596465950	50420831
Averaging	204605775027	5838956725	537926381	40272104
Luma	207067812232	5790118160	496829390	35546606
Decomposition (Max) ¹	-	-	-	-
Decomposition (Min)	206050242309	5817410748	516984118	37468228

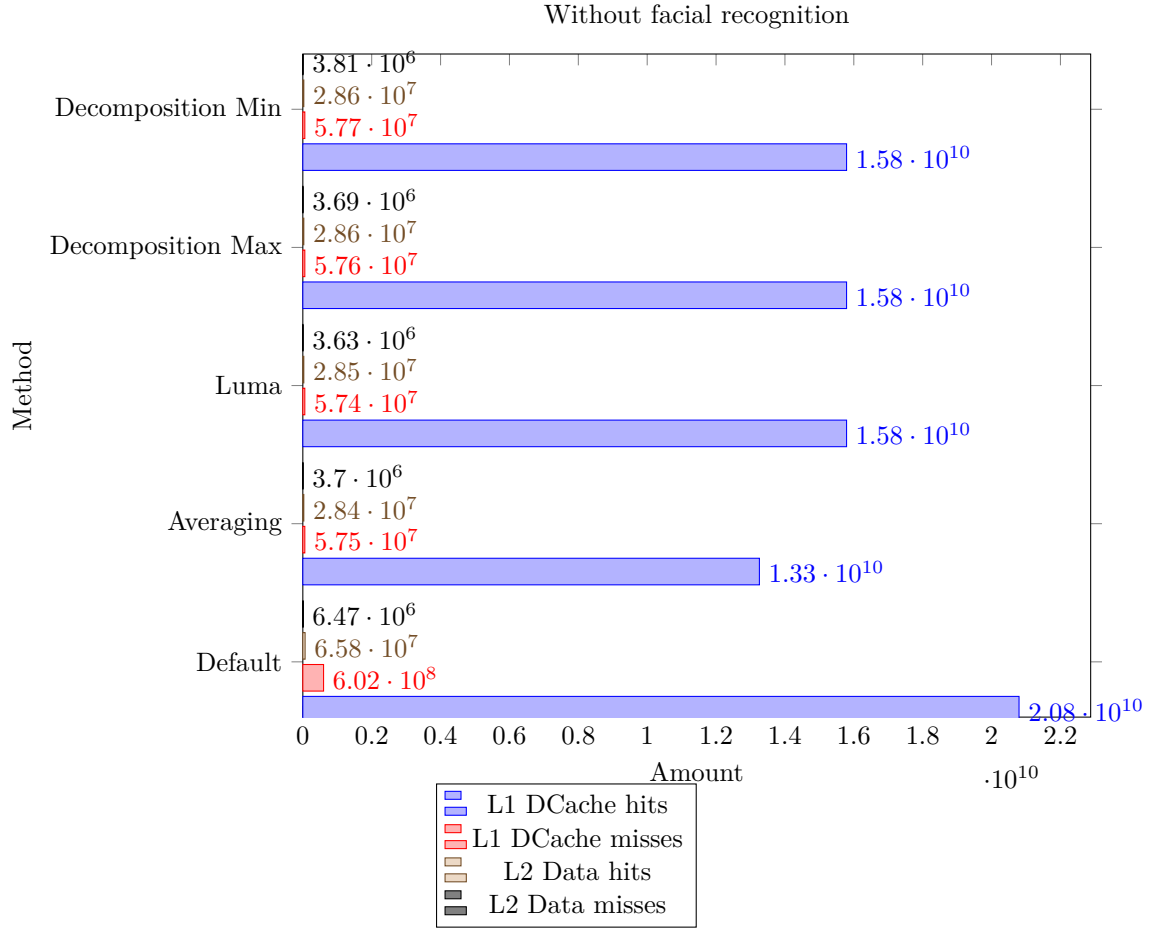
Without facial recognition:

Method	L1 DCache hits	L1 DCache misses	L2 Data hits	L2 Data misses
Default	20798177177	602197117	65785000	6469704
Averaging	13258521614	57533588	28390788	3701288
Luma	15789762340	57418201	28490336	3625812
Decomposition (Max)	15789802982	57631568	28642166	3686989
Decomposition (Min)	15789815633	57681407	28622821	3807790

¹**Decomposition (Max)** produced images so bad that these were unable to be used for facial recognition. This led to incorrect execution of the software, and therefore incorrect data.

4.2 Visualized





5 Processing

As predicted in our hypothesis the different methods are within margin of each other. Without facial recognition, it seems that *Default* has a bit more L1-DCache hits but also more L1-DCache misses.

According to our measurements it seems that both *Decomposition* and *Luma* perform almost identical.

There is one observation that can be made regarding data sizes and the cache. That it seems that the L1 DCache is used a lot more than the L2 cache. It seems that the cache lanes are big enough to fit a few pixels in, and it seems like the pre-fetcher recognizes that we are sequentially reading these pixels. One problem appears here, we can't see what is in the cache. Because the cache is transparent to the actual OS, it is impossible to see what is truly in the cache. This makes impossible to verify the pixels are being stored here, also it makes it impossible to see how many pixels are stored in the cache.

6 Conclusion

From these measurements we can't come to a clear conclusion about which is the better method. It seems that data from different aspects of these methods is needed to declare a winner of sorts. We can only really state that *Decomposition Max* is inadvisable as it produces bad images.

7 Evaluation

7.1 What went well

The project was a very cooperative experience for both parties involved. This cooperation went very well, everyone had their tasks and did them with equal amount of dedication. This made the project an overall pleasant experience.

7.2 What went wrong

We both had a vigorous start with the project, however it turned out to be very hard to keep up this speed. Eventually we started to have to divide our time to other projects that required our attention at the time. This made it so we had to do a lot of work near the deadline.

Also the original plan was to use OpenCL and SIMD for parallelism. Because of the state of the relevant tools, this turned out to be quite a bit harder than was originally thought.

7.3 What could be done differently

A better research on the more ambitious subjects should be done in order to avoid unnecessary delays.

8 Additional information

In this section there will be additional information about the system where the tests were run on.

8.1 LSCPU

```
Architecture:      x86_64
CPU op-mode(s):    32-bit, 64-bit
Byte Order:        Little Endian
CPU(s):            4
On-line CPU(s) list: 0-3
Thread(s) per core: 1
Core(s) per socket: 4
Socket(s):         1
NUMA node(s):      1
Vendor ID:         GenuineIntel
CPU family:        6
Model:            94
Model name:        Intel(R) Core(TM) i5-6600K CPU @ 3.50GHz
Stepping:          3
CPU MHz:           4100.012
CPU max MHz:       4100.0000
CPU min MHz:       800.0000
BogoMIPS:          7008.00
Virtualization:    VT-x
L1d cache:         32K
L1i cache:         32K
L2 cache:          256K
L3 cache:          6144K
NUMA node0 CPU(s): 0-3
Flags:             fpu vme de pse tsc msr pae mce cx8 apic sep mtrr
                   pge mca cmov pat pse36 clflush dts acpi mmx fxsr sse sse2 ss ht
                   tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc art arch_perfmon
                   pebs bts rep_good nopl xtopology nonstop_tsc cpuid aperfmperf
                   tsc_known_freq pni pclmulqdq dtes64 monitor ds_cpl vmx est tm2
                   ssse3 sdbg fma cx16 xtpr pdcm pcid sse4_1 sse4_2 x2apic movbe
                   popcnt aes xsave avx f16c rdrand lahf_lm abm 3dnowprefetch
                   cpuid_fault invpcid_single pti tpr_shadow vnmi flexpriority ept
                   vpid fsgsbase tsc_adjust bmi1 hle avx2 smep bmi2 erms invpcid rtm
                   mpx rdseed adx smap clflushopt intel_pt xsaveopt xsavec xgetbv1
                   xsaves dtherm ida arat pln pts hwp hwp_notify hwp_act_window
                   hwp_epp
```

References

- [1] Intel Ark product specifications <https://ark.intel.com/products/88191/Intel-Core-i5-6600K>